

Imports

```
In [ ]: import numpy as np
import pandas as pd
import matplotlib.pyplot as plt
import seaborn as sns
import plotly.express as px

from tqdm import tqdm
import math
import time

from sklearn.model_selection import train_test_split
from sklearn.model_selection import GridSearchCV
from sklearn.model_selection import validation_curve
from sklearn.model_selection import learning_curve
from sklearn.preprocessing import StandardScaler

from sklearn.tree import DecisionTreeClassifier
from sklearn.ensemble import GradientBoostingClassifier, AdaBoostClassifier
from sklearn.neighbors import KNeighborsClassifier
from sklearn.svm import SVC
from sklearn.neural_network import MLPClassifier

from sklearn.metrics import classification_report, confusion_matrix
from sklearn.metrics import accuracy_score

from pathlib import Path
import joblib
```

User Config

```
In [ ]: class Config:
    def __init__(self):
        self.dataset = 'Data/diabetes.csv'
        self.dataset_name = "Diabetes"
        self.features = ['Pregnancies',
                        'PlasmaGlucose',
                        'DiastolicBloodPressure',
                        'TricepsThickness',
                        'SerumInsulin',
                        'BMI',
                        'DiabetesPedigree',
                        'Age'
                        ]
        self.label = 'Diabetic'

#         self.dataset = 'Data/penguins.csv'
#         self.features = ['CulmenLength',
#                         'CulmenDepth',
#                         'FlipperLength',
#                         'BodyMass'
#                         ]
```

```
#         self.label = 'Species'

#         self.dataset = 'Data/breast-cancer-wisconsin.csv'
#         self.dataset_name = "Breast Cancer"
#         self.features = ["Clump Thickness",
#                           "Uniformity of Cell Size",
#                           "Uniformity of Cell Shape",
#                           "Marginal Adhesion",
#                           "Single Epithelial Cell Size",
#                           "Bare Nuclei",
#                           "Bland Chromatin",
#                           "Normal Nucleoli",
#                           "Mitoses"
#                           ]
#         self.label = 'Class'

self.random_seed = 101

config = Config()
```

Read File

```
In [ ]: # Load the training dataset
df = pd.read_csv(config.dataset)
df.head()
```

EDA

```
In [ ]: # Understand number of rows and columns
df.shape
```

```
In [ ]: # Understand what your columns are and data types
df.info()
```

```
In [ ]: # Understand the descriptive statistics of your numerical data
df.describe()
```

```
In [ ]: # Check for null values
df.isna().sum() / len(df)*100
# sns.heatmap(df.isnull(), yticklabels=False, cbar=False, cmap='viridis')
```

```
In [ ]: # notice imbalanced data set
df[config.label].value_counts()
```

```
In [ ]: # benign = df[df[config.label]==2]
# benign_sample = benign.sample(n=239, replace=False, random_state=config.random_seed,
# benign_sample

# malig = df[df[config.label]==4]
# malig

# df = pd.concat([benign_sample, malig], ignore_index=True)
# df['Class'] = [0 if x==2 else 1 for x in df['Class']] ]
```

```

In [ ]: # peng0 = df[df[config.Label]==0]
# peng0_sample = peng0.sample(n=68, replace=False, random_state=config.random_seed, ig
# peng0_sample

# peng1 = df[df[config.Label]==1]
# peng1_sample = peng1.sample(n=68, replace=False, random_state=config.random_seed, ig
# peng1_sample

# peng2 = df[df[config.Label]==2]
# peng2

# df = pd.concat([peng0_sample, peng1_sample, peng2], ignore_index=True)

In [ ]: non_diabetic = df[df['Diabetic']==0]
non_diabetic_sample = non_diabetic.sample(frac=0.5, replace=False, random_state=config
non_diabetic_sample

diabetic = df[df['Diabetic']==1]
diabetic

df = pd.concat([non_diabetic_sample, diabetic], ignore_index=True)

In [ ]: df[config.label].value_counts()

In [ ]: for col in config.features:
df.boxplot(column=col, by=config.label, figsize=(6,4))
# plt.title(col)
plt.show()

In [ ]: for col in config.features:
# fg = sns.FacetGrid(df, col=config.Label)
# fg.map(plt.hist, col)
plt.figure(figsize=(6,2))
sns.histplot(data=df, x=col, hue=config.label)
plt.show()

In [ ]: sns.pairplot(df[config.features + [config.label]], hue=config.label, diag_kind="hist")

```

Train Test Split

```

In [ ]: X = df[config.features].values

y = df[config.label].values

In [ ]: # Split data 70%-30% into training set and test set
X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.15, random_state

print(f"Training Cases: {len(X_train)} \nTest Cases: {len(X_test)}")

In [ ]:

In [ ]:

```

Modelling

```
In [ ]: # Decision trees with some form of pruning
        # Neural networks
        # Boosting
        # Support Vector Machines
        # k-nearest neighbors
```

```
In [ ]: # train a Decision Tree Classifier model on the training set
model = DecisionTreeClassifier(criterion='gini',
                               splitter='best',
                               max_depth=None,
                               min_samples_split=2,
                               min_samples_leaf=1,
                               min_weight_fraction_leaf=0.0,
                               max_features=None,
                               random_state=None, #config.random_seed,
                               max_leaf_nodes=None,
                               min_impurity_decrease=0.0,
                               class_weight=None,
                               ccp_alpha=0.0,
                               )

# model = AdaBoostClassifier(base_estimator=None,
#                             #
#                             n_estimators=50,
#                             #
#                             learning_rate=1.0,
#                             #
#                             algorithm='SAMME.R',
#                             #
#                             random_state=None
#                             )

# model = KNeighborsClassifier(n_neighbors=5,
#                              #
#                              weights='uniform',
#                              #
#                              algorithm='auto',
#                              #
#                              leaf_size=30,
#                              #
#                              p=2,
#                              #
#                              metric='minkowski',
#                              #
#                              metric_params=None,
#                              #
#                              n_jobs=-1
#                              )

model = SVC(C=1.0,
            kernel='rbf',
            degree=3,
            gamma='scale',
            coef0=0.0,
            shrinking=True,
            probability=False,
            tol=0.001,
            cache_size=200,
            class_weight=None,
            verbose=False,
            max_iter=-1,
            decision_function_shape='ovr',
            break_ties=False,
            random_state=None
            )

# model = MLPClassifier(hidden_layer_sizes=(100,),
```

```
#         activation='relu',
#         solver='adam',
#         alpha=0.0001,
#         batch_size='auto',
#         learning_rate='constant',
#         learning_rate_init=0.001,
#         power_t=0.5,
#         max_iter=200,
#         shuffle=True,
#         random_state=None,
#         tol=0.0001,
#         verbose=False,
#         warm_start=False,
#         momentum=0.9,
#         nesterovs_momentum=True,
#         early_stopping=False,
#         validation_fraction=0.1,
#         beta_1=0.9,
#         beta_2=0.999,
#         epsilon=1e-08,
#         n_iter_no_change=10,
#         max_fun=15000
#     )
```

In []:

Preprocessing

In []:

```
# # Train the model
# from sklearn.compose import ColumnTransformer
# from sklearn.pipeline import Pipeline
# from sklearn.preprocessing import StandardScaler

# # Define preprocessing for numeric columns (normalize them so they're on the same scale)
# # numeric_features = [0,1,2,3,4,5,6]
# # numeric_transformer = Pipeline(steps=[
# #     ('scaler', StandardScaler())])

# # Define preprocessing for categorical features (encode the Age column)
# # categorical_features = [7]
# # categorical_transformer = Pipeline(steps=[
# #     ('onehot', OneHotEncoder(handle_unknown='ignore'))])

# # Combine preprocessing steps
# # preprocessor = ColumnTransformer(
# #     transformers=[
# #         ('num', numeric_transformer, numeric_features),
# #         ('cat', categorical_transformer, categorical_features)])

# # Create preprocessing and training pipeline
# # pipeline = Pipeline(steps=[('preprocessor', preprocessor),
# #                             ('logregressor', model)])

# # Create preprocessing and training pipeline
# # pipeline = Pipeline(steps=[('scaler', StandardScaler()),
# #                             ('knn', model)])
```

```
# fit the pipeline to train a logistic regression model on the training set
# model = pipeline.fit(X_train, y_train)
# print (model)
```

```
In [ ]: if isinstance(model, KNeighborsClassifier) or isinstance(model, SVC) or isinstance(model, LogisticRegression):
        scaler = StandardScaler()

        X_train = scaler.fit_transform(X_train)

        X_test = scaler.transform(X_test)

        X = scaler.fit_transform(X) # For timing and validation curve experiments below
```

Single Model Fit and Predict

```
In [ ]: model.fit(X_train, y_train)

        predictions = model.predict(X_test)
```

Single Model Evaluation

```
In [ ]: print('Accuracy: ', accuracy_score(y_test, predictions))
        print(confusion_matrix(y_test, predictions))

        print(classification_report(y_test, predictions))
```

```
In [ ]: if isinstance(model, DecisionTreeClassifier) or isinstance(model, AdaBoostClassifier):
        print(model.feature_importances_)
```

```
In [ ]: config.features
```

```
In [ ]:
```

```
In [ ]: # from sklearn.metrics import roc_curve
        # from sklearn.metrics import confusion_matrix

        # y_scores = model.predict_proba(X_test)

        # calculate ROC curve
        # fpr, tpr, thresholds = roc_curve(y_test, y_scores[:,1])

        # plot ROC curve
        # fig = plt.figure(figsize=(6, 6))
        # Plot the diagonal 50% line
        # plt.plot([0, 1], [0, 1], 'k--')
        # Plot the FPR and TPR achieved by our model
        # plt.plot(fpr, tpr)
        # plt.xlabel('False Positive Rate')
        # plt.ylabel('True Positive Rate')
        # plt.title('ROC Curve')
        # plt.show()
```

```
In [ ]: # from sklearn.metrics import roc_auc_score

# auc = roc_auc_score(y_test,y_scores[:,1])
# print('AUC: ' + str(auc))
```

Experiments

Gridsearch Experiment

```
In [ ]: %%time

if isinstance(model, DecisionTreeClassifier):
    param_grid = {
        'criterion' : ["gini", "entropy", "log_loss"],
        'max_depth' : [int(x) for x in np.linspace(1, 20, num = 20)],
#         'min_samples_split' : [int(x) for x in np.linspace(2, 50, num = 10)],
#         'min_samples_leaf' : [int(x) for x in np.linspace(1, 20, num = 10)],
#         'max_features' : [int(x) for x in np.linspace(2, X.shape[1], num = X.shape[1]
    }
elif isinstance(model, AdaBoostClassifier):
    param_grid = {
        "n_estimators": [int(x) for x in np.linspace(1, 100, num = 20)],
        "learning_rate": [float(x) for x in np.linspace(0.1, 2.0, 20)],
        "algorithm": ['SAMME', 'SAMME.R'],
    }
elif isinstance(model, KNeighborsClassifier):
    param_grid = {
        "n_neighbors": [int(x) for x in np.linspace(1, 50, num = 50)],
#         "algorithm": ['auto', 'ball_tree', 'kd_tree', 'brute'] ,
        "p": [int(x) for x in np.linspace(1, 3, num = 3)],
    }
elif isinstance(model, SVC):
    param_grid = {
        "C": [float(x) for x in np.linspace(0.1, 10.1, 21)],
        "kernel": ["linear", "poly", "rbf", "sigmoid"],
#         degree=3,
#         gamma='scale',
#         coef0=0.0,
#         shrinking=True,
#         probability=False,
#         tol=0.001,
#         cache_size=200,
#         class_weight=None,
#         verbose=False,
#         max_iter=-1,
#         decision_function_shape='ovr',
#         break_ties=False,
#         random_state=None
    }
else: #MLPClassifier()
    param_grid = {
        "hidden_layer_sizes": [(int(x),) for x in np.linspace(1, 150, num = 16)] + [(1
        'max_iter': [500,1000,1500,2000],
        'alpha': 10.0 ** -np.arange(1, 7),
```

```
#         activation='relu',
#         'solver':['lbfgs', 'sgd', 'adam'],
#         alpha=0.0001,
#         batch_size='auto',
#         'learning_rate':['constant', "invscaling", "adaptive"]
#         learning_rate_init=0.001,
#         power_t=0.5,
#         max_iter=200,
#         shuffle=True,
#         random_state=None,
#         tol=0.0001,
#         verbose=False,
#         warm_start=False,
#         momentum=0.9,
#         nesterovs_momentum=True,
#         early_stopping=True,
#         validation_fraction=0.1,
#         beta_1=0.9,
#         beta_2=0.999,
#         epsilon=1e-08,
#         n_iter_no_change=100,
#         max_fun=15000
#     }

grid_search = GridSearchCV(model, param_grid, cv=5, verbose=2, n_jobs=-1)

grid_search.fit(X, y)
```

```
In [ ]: # %%time
# param_grid = {
#     "n_neighbors": [int(x) for x in np.linspace(, 1000, num = 20)],
#     "weights": ['uniform', 'distance'] ,
#     "p": [int(x) for x in np.linspace(1, 12, num = 12)],
# }

# pipeline = Pipeline(steps=[('scaler', StandardScaler()),
#                             ('grid', GridSearchCV(model, param_grid, cv=5, verbose=2,
# pipeline.fit(X_train, y_train)
```

```
In [ ]: grid_search_experiment_df = pd.DataFrame(grid_search.cv_results_)
grid_search_experiment_df.head()
```

```
In [ ]: grid_search.best_params_
```

```
In [ ]: grid_search.best_score_
```

```
In [ ]: if isinstance(model, DecisionTreeClassifier) or isinstance(model, AdaBoostClassifier):
#     print(grid_search.best_estimator_.feature_importances_)
dfz = pd.DataFrame(zip(config.features, grid_search.best_estimator_.feature_importances_))
dfz.sort_values(by="Feature Importance", ascending=False, inplace=True)
dfz.plot(kind='bar', x='Feature', y='Feature Importance', title=f"Feature Importance")
```

```
In [ ]: config.features
```

```
In [ ]:
```

```
In [ ]: # grid_search_experiment_df[(grid_search_experiment_df['param_criterion'] == 'entropy
```



```
# # & (grid_search_experiment_df['param_min_samples_leaf'] =
# # & (grid_search_experiment_df['param_min_samples_split']
# # & (grid_search_experiment_df['param_max_depth'] == 9)
# ].plot("param_max_depth", "mean_test_score", style="o-")
```

```
In [ ]: fig, ax = plt.subplots()
grid_search_experiment_df.pivot_table(values="mean_test_score", index="param_max_depth",
# sns.lineplot("param_max_depth", "mean_test_score", hue="param_criterion", data=grid_
ax.set_xticks(param_grid['max_depth'])
ax.set_ylabel("Score")
ax.set_title(f"Accuracy Vs. max_depth - {config.dataset_name}")
ax.grid()
```

```
In [ ]: # same plot using Seaborn
# fig = px.line(grid_search_experiment_df, x="param_max_depth", y="mean_test_score", c
# fig.show()
```

Timing Experiment

```
In [ ]: # splits = [0.1, 0.2, 0.3, 0.4, 0.5, 0.6, 0.7, 0.8, 0.9, 0.99]
# reps = 5 # number of times to run each experiment for mean calculation

# fit_times = np.zeros(shape=(len(splits), reps))
# pred_times = np.zeros(shape=(len(splits), reps))
# test accuracies = np.zeros(shape=(len(splits), reps))
# train accuracies = np.zeros(shape=(len(splits), reps))

# for i, split in enumerate(tqdm(splits)):
#     for j in range(reps):
#         X_train, X_test, y_train, y_test = train_test_split(X, y, train_size=split,)

#         # Get time to fit model
#         start_time = time.time()
#         grid_search.best_estimator_.fit(X_train, y_train)
#         end_time = time.time()
#         wall_time = end_time - start_time
#         fit_times[i, j] = wall_time

#         # Get time to make test predictions
#         start_time = time.time()
#         test_predictions = grid_search.best_estimator_.predict(X_test)
#         end_time = time.time()
#         wall_time = end_time - start_time
#         pred_times[i, j] = wall_time

#         #train_predictions = model.predict(X_train)

#         train_accuracies[i, j] = grid_search.best_estimator_.score(X_train, y_train)
#         test_accuracies[i, j] = grid_search.best_estimator_.score(X_test, y_test) #c

# # Calculate the average of reps
# mean_train_accuracies = np.mean(train_accuracies, axis=1)
# mean_test_accuracies = np.mean(test_accuracies, axis=1)
# mean_fit_times = np.mean(fit_times, axis=1)
# mean_pred_times = np.mean(pred_times, axis=1)

# # Store results in df
```

```
# time_experiment_df = pd.DataFrame({"train_size": splits,
#                                   "mean_fit_times": mean_fit_times,
#                                   "mean_pred_times": mean_pred_times,
#                                   "mean_train_accuracies": mean_train_accuracies,
#                                   "mean_test_accuracies": mean_test_accuracies})
# time_experiment_df
```

```
In [ ]: # time_experiment_df['mean_train_loss'] = 1 - time_experiment_df['mean_train_accuracies']
# time_experiment_df['mean_test_loss'] = 1 - time_experiment_df['mean_test_accuracies']

# time_experiment_df.plot(x="train_size", y=["mean_train_loss", "mean_test_loss"])
# # time_experiment_df.plot(x="train_size", y=["mean_test_accuracies", "mean_train_accuracies"])
```

Learning Curve

```
In [ ]: %%time

# Code shamelessly copied from: https://scikit-learn.org/stable/auto_examples/model_selection/learning_curve.html

# train_sizes, train_scores, test_scores, fit_times, score_times = learning_curve(gri
train_sizes, train_scores, test_scores, fit_times, score_times = learning_curve(tuned
train_scores_mean = np.mean(train_scores, axis=1)
train_scores_std = np.std(train_scores, axis=1)
test_scores_mean = np.mean(test_scores, axis=1)
test_scores_std = np.std(test_scores, axis=1)
fit_times_mean = np.mean(fit_times, axis=1)
fit_times_std = np.std(fit_times, axis=1)
score_times_mean = np.mean(score_times, axis=1)
score_times_std = np.std(score_times, axis=1)

fig, axes = plt.subplots(nrows=1, ncols=2, figsize=(20, 5))

# Plot Learning curve
axes[0].grid()
axes[0].fill_between(
    train_sizes,
    train_scores_mean - train_scores_std,
    train_scores_mean + train_scores_std,
    alpha=0.1,
    color="r",
)
axes[0].fill_between(
    train_sizes,
    test_scores_mean - test_scores_std,
    test_scores_mean + test_scores_std,
    alpha=0.1,
    color="g",
)
axes[0].plot(
    train_sizes, train_scores_mean, "o-", color="r", label="Training score"
)
axes[0].plot(
    train_sizes, test_scores_mean, "o-", color="g", label="Cross-validation score"
)
axes[0].legend(loc="best")
axes[0].set_xlabel("Training examples")
axes[0].set_ylabel("Score")
axes[0].set_title(f"Learning Curves - {config.dataset_name}")
```

```

# # Plot n_samples vs fit_times
# axes[1].grid()
# axes[1].plot(train_sizes, fit_times_mean, "o-")
# axes[1].fill_between(
#     train_sizes,
#     fit_times_mean - fit_times_std,
#     fit_times_mean + fit_times_std,
#     alpha=0.1,
# )
# axes[1].set_xlabel("Training examples")
# axes[1].set_ylabel("fit_times")
# axes[1].set_title(f"Scalability of the model - {config.dataset_name}")

```

In []:

Loss Curve for MLP

In []:

```

%%time

fig, ax1 = plt.subplots(nrows=1, ncols=1, figsize=(10, 5))
ax1.plot(grid_search.best_estimator_.loss_curve_, label="best_estimator10e-3")

mlp_pt1 = MLPClassifier(alpha=1e-05,
                        hidden_layer_sizes=(130,),
                        max_iter=1500,
                        early_stopping=True,
                        n_iter_no_change=100,
                        learning_rate_init = 0.1)
mlp_pt01 = MLPClassifier(alpha=1e-05,
                        hidden_layer_sizes=(130,),
                        max_iter=1500,
                        early_stopping=True,
                        n_iter_no_change=100,
                        learning_rate_init = 0.01)
# mlp_pt001 = MLPClassifier(alpha=1e-05,
#                           hidden_layer_sizes=(130,),
#                           max_iter=1500,
#                           early_stopping=True,
#                           n_iter_no_change=100,
#                           learning_rate_init = 0.001)
mlp_pt0001 = MLPClassifier(alpha=1e-05,
                          hidden_layer_sizes=(130,),
                          max_iter=1500,
                          early_stopping=True,
                          n_iter_no_change=100,
                          learning_rate_init = 0.0001)

mlp_pt1.fit(X_train, y_train)
mlp_pt01.fit(X_train, y_train)
# mlp_pt001.fit(X_train, y_train)
mlp_pt0001.fit(X_train, y_train)

ax1.plot(mlp_pt1.loss_curve_, label="learning_rate10e-1", alpha=0.5)
ax1.plot(mlp_pt01.loss_curve_, label="learning_rate10e-2", alpha=0.5)
# ax1.plot(mlp_pt001.loss_curve_, label="learning_rate10e-1")
ax1.plot(mlp_pt0001.loss_curve_, label="learning_rate10e-4", alpha=0.5)
ax1.legend()

```

```
plt.title(f"Loss Curve - {config.dataset_name}")
plt.xlabel(f"Iterations")
plt.ylabel("Loss")
```

Validation Curves

```
In [ ]: tuned_model = joblib.load("./Models/hypertuned_MLPClassifier().pkl")
```

```
In [ ]: %%time
# Code shamelessly copied from: https://scikit-learn.org/stable/auto_examples/model_selection/plot_validation_curve.html

# DecisionTree
param_name = "max_depth"
param_range = [int(x) for x in np.linspace(1, 20, num = 20)]
# param_name = "criterion"
# param_range = ["gini", "entropy", "log_loss"]

# AdaBoost
param_name = "learning_rate"
param_range = [float(x) for x in np.linspace(0.1, 1.8, 18)]
param_name = "n_estimators"
param_range = [int(x) for x in np.linspace(1, 200, num = 20)]

# KNN
param_name = "n_neighbors"
param_range = [int(x) for x in np.linspace(1, 30, 30)]
param_name = "p"
param_range = [int(x) for x in np.linspace(1, 12, 12)]

# SVM
param_name = "C"
param_range = [float(x) for x in np.linspace(0.1, 10.1, 21)]
# param_name = "kernel"
# param_range = ["linear", "poly", "rbf", "sigmoid"]

# # NN
param_name = "hidden_layer_sizes"
param_range = [(int(x),) for x in np.linspace(1, 150, num = 16)]
# param_range = [(30,), (30,30), (30,30,30), (30,30,30,30)]
param_range = [(130,), (130,130), (130,130,130), (130,130,130,130)]

# param_name = 'activation'
# param_range = ["identity", "logistic", "tanh", "relu"]
param_name = "max_iter"
param_range = list(np.arange(100, 2000, 100))
# param_name = "solver"
# param_range = ["lbfgs", "sgd", "adam"]
# param_name = "learning_rate"
# param_range = ["constant", "invscaling", "adaptive"]

train_scores, test_scores = validation_curve(
    tuned_model, #grid_search.best_estimator_,
    X,
    y,
    param_name=param_name,
    param_range=param_range,
    scoring="accuracy",
```

```

    n_jobs=-1,
)
train_scores_mean = np.mean(train_scores, axis=1)
train_scores_std = np.std(train_scores, axis=1)
test_scores_mean = np.mean(test_scores, axis=1)
test_scores_std = np.std(test_scores, axis=1)

# specifically to for hidden_layer_sizes of NN
if param_name == "hidden_layer_sizes":
    param_range = [str(x) for x in param_range]

plt.figure(figsize=(6,3))
lw = 2
plt.plot(
    param_range, train_scores_mean, label="Training score", color="darkorange", lw=lw,
)
plt.fill_between(
    param_range,
    train_scores_mean - train_scores_std,
    train_scores_mean + train_scores_std,
    alpha=0.2,
    color="darkorange",
    lw=lw,
)
plt.plot(
    param_range, test_scores_mean, label="Cross-validation score", color="navy", lw=lw,
)
plt.fill_between(
    param_range,
    test_scores_mean - test_scores_std,
    test_scores_mean + test_scores_std,
    alpha=0.2,
    color="navy",
    lw=lw,
)
plt.title(f"Validation Curve with {model} - {config.dataset_name}")
plt.xlabel(f"{param_name}")
plt.ylabel("Score")
# plt.xscale('log')
# plt.ylim(0.7, 1.1)
plt.legend(loc="best")
plt.grid()
plt.xticks(param_range)
plt.xticks(rotation=45)
plt.show()

```

Save Model

```

In [ ]: ## Save the model as a pickle file
# filename = f'./Models/{config.dataset_name}/hypertuned_{model}.pkl'
# joblib.dump(grid_search.best_estimator_, filename)

```

Comparison Between Models

```

In [ ]: %%time

```

```

# Code shamelessly copied from: https://scikit-learn.org/stable/auto_examples/model_se

# assign directory
directory = './Models/'

# iterate over files in that directory
tuned_models = Path(directory).glob('hypertuned_*')

for tuned_model in tuned_models:

    tuned_model = joblib.load(tuned_model)
    if isinstance(tuned_model, DecisionTreeClassifier):
        label = "DecisionTreeClassifier"
    elif isinstance(tuned_model, AdaBoostClassifier):
        label = "AdaBoostClassifier"
    elif isinstance(tuned_model, KNeighborsClassifier):
        label = "KNeighborsClassifier"
    elif isinstance(tuned_model, SVC):
        label = "SVC"
    else:
        label = "MLPClassifier"
#     continue

    train_sizes, train_scores, test_scores, fit_times, score_times = learning_curve(t
    train_scores_mean = np.mean(train_scores, axis=1)
    train_scores_std = np.std(train_scores, axis=1)
    test_scores_mean = np.mean(test_scores, axis=1)
    test_scores_std = np.std(test_scores, axis=1)
    fit_times_mean = np.mean(fit_times, axis=1)
    fit_times_std = np.std(fit_times, axis=1)
    score_times_mean = np.mean(score_times, axis=1)
    score_times_std = np.std(score_times, axis=1)

    fig, ax = plt.subplots(nrows=1, ncols=1, figsize=(4, 2))

    # Plot Learning curve
    ax.grid()
    ax.fill_between(
        train_sizes,
        train_scores_mean - train_scores_std,
        train_scores_mean + train_scores_std,
        alpha=0.1,
        color="r",
    )
    ax.fill_between(
        train_sizes,
        test_scores_mean - test_scores_std,
        test_scores_mean + test_scores_std,
        alpha=0.1,
        color="g",
    )
    ax.plot(
        train_sizes, train_scores_mean, "o-", color="r", label="Training score"
    )
    ax.plot(
        train_sizes, test_scores_mean, "o-", color="g", label="Cross-validation score"
    )
    ax.legend(loc="best")

```

```
ax.set_xlabel("Training examples")
ax.set_ylabel("Score")
ax.set_title(f"{label} Learning Curves")
ax.set_ylim(0.85, 1.03)
```

```
In [ ]: %%time
# Create a graph of combined fit times
# Code shamelessly copied from: https://scikit-learn.org/stable/auto_examples/model_se

# assign directory
directory = './Models/'

# iterate over files in that directory
tuned_models = Path(directory).glob('hypertuned_*')

fig, ax = plt.subplots(nrows=1, ncols=2, figsize=(23, 5))

for tuned_model in tuned_models:

    tuned_model = joblib.load(tuned_model)

    if isinstance(tuned_model, DecisionTreeClassifier):
        label = "DecisionTreeClassifier"
        color = "orange"
    elif isinstance(tuned_model, AdaBoostClassifier):
        label = "AdaBoostClassifier"
        color = "tab:blue"
    elif isinstance(tuned_model, KNeighborsClassifier):
        label = "KNeighborsClassifier"
        color = "green"
    elif isinstance(tuned_model, SVC):
        label = "SVC"
        color = "red"
    else:
        label = "MLPClassifier"
        color = "tab:purple"
    continue

    train_sizes, train_scores, test_scores, fit_times, score_times = learning_curve(t
    fit_times, score_times = fit_times*1000, score_times*1000
    train_scores_mean = np.mean(train_scores, axis=1)
    train_scores_std = np.std(train_scores, axis=1)
    test_scores_mean = np.mean(test_scores, axis=1)
    test_scores_std = np.std(test_scores, axis=1)
    fit_times_mean = np.mean(fit_times, axis=1)
    fit_times_std = np.std(fit_times, axis=1)
    score_times_mean = np.mean(score_times, axis=1)
    score_times_std = np.std(score_times, axis=1)

    if isinstance(tuned_model, KNeighborsClassifier):
        score_times_mean[0] = np.min(score_times[0])

    # Plot n_samples vs fit_times
    ax[0].grid()
    ax[0].plot(train_sizes, fit_times_mean, "o-", color=color, label=label)
    ax[0].fill_between(
        train_sizes,
        fit_times_mean - fit_times_std,
        fit_times_mean + fit_times_std,
```

```

        alpha=0.1,
        color=color
    )

    # Plot n_samples vs fit_times
    ax[1].grid()
    ax[1].plot(train_sizes, score_times_mean, "o-", color=color, label=label)
    # ax[1].fill_between(
    #     train_sizes,
    #     score_times_mean - score_times_std,
    #     score_times_mean + score_times_std,
    #     alpha=0.1,
    # )

    ax[0].set_xlabel("Training examples", fontsize=18)
    ax[0].set_ylabel("Training fit_times (ms)", fontsize=18)
    ax[0].set_title("Scalability of the model - Training", fontsize=25)
    ax[0].tick_params(axis='both', which='major', labelsize=15)
    ax[0].legend(fontsize=15)
    ax[0].grid()

    ax[1].set_xlabel("Training examples", fontsize=18)
    ax[1].set_ylabel("Predict Times (ms)", fontsize=18)
    ax[1].set_title("Scalability of the model - Predicting", fontsize=25)
    ax[1].tick_params(axis='both', which='major', labelsize=15)
    ax[1].legend(fontsize=15)
    ax[1].grid()

```

Scratch Work

```

In [ ]: # #Visualizing KNN Decision Boundary

# from matplotlib.colors import ListedColormap
# from sklearn.inspection import DecisionBoundaryDisplay

# n_neighbors = 9

# X = df[config.features].iloc[:, [0]+[-1]]

# # Create color maps
# cmap_light = ListedColormap(["orange", "cyan"])
# cmap_bold = ["darkorange", "c"]

# for weights in ["uniform", "distance"]:
#     # we create an instance of Neighbours Classifier and fit the data.
#     clf = KNeighborsClassifier(n_neighbors, weights=weights)
#     clf.fit(X, y)

#     _, ax = plt.subplots()
#     DecisionBoundaryDisplay.from_estimator(
#         clf,
#         X,
#         cmap=cmap_light,
#         ax=ax,
#         response_method="predict",
#         plot_method="pcolormesh",
#         xlabel=X.columns[0],

```



```

#         ylabel=X.columns[1],
#         shading="auto",
#     )

#     # Plot also the training points
#     sns.scatterplot(
#         x=X.iloc[:, 0],
#         y=X.iloc[:, 1],
#         hue=y,
#         palette=cmap_bold,
#         alpha=1.0,
#         edgecolor="black",
#     )
#     plt.title(
#         "KNN Classifier (k = %i, weights = '%s')" % (n_neighbors, weights)
#     )

# plt.show()

```

```

In [ ]: # from sklearn import svm
# from sklearn.inspection import DecisionBoundaryDisplay

# X = df[config.features].iloc[:, [0]+[-1]]

# # we create an instance of SVM and fit out data. We do not scale our
# # data since we want to plot the support vectors
# C = 1.0 # SVM regularization parameter
# models = (
#     svm.SVC(kernel="linear", C=C),
#     svm.LinearSVC(C=C, max_iter=10000),
#     svm.SVC(kernel="rbf", gamma=0.7, C=C),
#     svm.SVC(kernel="poly", degree=3, gamma="auto", C=C),
# )
# models = (clf.fit(X, y) for clf in models)

# # title for the plots
# titles = (
#     "SVC with linear kernel",
#     "LinearSVC (linear kernel)",
#     "SVC with RBF kernel",
#     "SVC with polynomial (degree 3) kernel",
# )

# # Set-up 2x2 grid for plotting.
# fig, sub = plt.subplots(2, 2)
# plt.subplots_adjust(wspace=0.4, hspace=0.4)

# X0, X1 = X.iloc[:, 0], X.iloc[:, 1]

# for clf, title, ax in zip(models, titles, sub.flatten()):
#     disp = DecisionBoundaryDisplay.from_estimator(
#         clf,
#         X,
#         response_method="predict",
#         cmap=plt.cm.coolwarm,
#         alpha=0.8,
#         ax=ax,
#         xlabel=X.columns[0],
#         ylabel=X.columns[1],

```

```
# )
# ax.scatter(X0, X1, c=y, cmap=plt.cm.coolwarm, s=20, edgecolors="k")
# ax.set_xticks(())
# ax.set_yticks(())
# ax.set_title(title)

# plt.show()
```

In []: